

Next-Generation Database Interfaces: A Survey of LLM-based Text-to-SQL

Zijin Hong¹, Zheng Yuan², Qinggang Zhang², Hao Chen²,
Junnan Dong², Feiran Huang¹, and Xiao Huang^{2*}

¹Jinan University, Guangzhou, China

²The Hong Kong Polytechnic University, Hong Kong SAR, China

hongzijin@stu2020.jnu.edu.cn, yzheng.yuan@connect.polyu.hk, qinggangg.zhang@connect.polyu.hk,
sundaychenhao@gmail.com, hanson.dong@connect.polyu.hk, huangfr@jnu.edu.cn, xiaohuang@comp.polyu.edu.hk

Abstract—Generating accurate SQL according to natural language questions (text-to-SQL) is a long-standing problem since it is challenging in user question understanding, database schema comprehension, and SQL generation. Conventional text-to-SQL systems include human engineering and deep neural networks. Subsequently, pre-trained language models (PLMs) have been developed and utilized for text-to-SQL tasks, achieving promising performance. As modern databases become more complex and corresponding user questions more challenging, PLMs with limited comprehension capabilities can lead to incorrect SQL generation. This necessitates more sophisticated and tailored optimization methods, which, in turn, restricts the applications of PLM-based systems. Most recently, large language models (LLMs) have demonstrated significant abilities in natural language understanding as the model scale remains increasing. Therefore, integrating the LLM-based implementation can bring unique opportunities, challenges, and solutions to text-to-SQL research. In this survey, we present a comprehensive review of LLM-based text-to-SQL. Specifically, we propose a brief overview of the current challenges and the evolutionary process of text-to-SQL. Then, we provide a detailed introduction to the datasets and metrics designed to evaluate text-to-SQL systems. After that, we present a systematic analysis of recent advances in LLM-based text-to-SQL. Finally, we discuss the remaining challenges in this field and propose expectations for future directions.

Index Terms—Text-to-SQL, Large Language Models, Database, Natural Language Processing

I. INTRODUCTION

TEXT-TO-SQL is a long-standing task in natural language processing research. It aims to convert (translate) natural language questions into database-executable SQL queries. Fig. 1 provides an example of the large language model-based (LLM-based) text-to-SQL; given a user question “*What cartoons were written by Joseph Kuhr?*”, the LLMs take the question and its corresponding database schema as the input and generate an SQL query as the output, which can be executed in the database to retrieve the content “*Batman Series*” for answering the user question. The above system builds a natural language interface to the database (NLIDB) with LLMs. Since SQL remains one of the most widely used programming languages, with over half (51.52%) of professional developers using SQL in their work, but only around a third (35.29%)

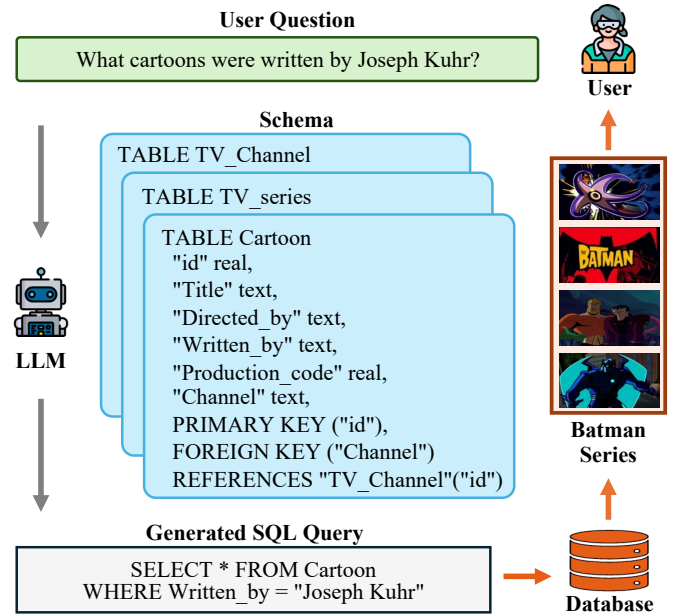


Fig. 1: An example for LLM-based text-to-SQL selected from Spider. The user proposes a question, “*What cartoons were written by Joseph Kuhr?*” the LLM takes the question and the schema of its corresponding database as the input, then generates a SQL query as the output. The SQL query can be executed in the database and retrieve a content “*Batman Series*” to answer the user question.

of those systematically trained¹, the NLIDB enables non-skilled users to access structured databases like professional database engineers [1, 2] and also accelerates human-computer interaction [3]. Furthermore, amid the research hotspot of LLMs, text-to-SQL can provide a solution to the prevalent hallucination [4, 5] issue by incorporating realistic content from the database to fill the knowledge gaps of LLMs [6]. The significant value and potential for text-to-SQL have triggered a range of studies on its incorporation and optimization with LLMs [7–9]; consequently, LLM-based text-to-SQL remains a highly discussed research field within the NLP and database

* Corresponding author.

¹<https://survey.stackoverflow.co/2023>




Fig. 2: A sketch of the evolutionary process for text-to-SQL research from the perspective of implementation paradigm. Each stage is presented with some implementation techniques and representative works. The timestamps for the stages are not exactly accurate; we set each timestamp according to the release time of the representative works of each paradigm, with a margin of error of about one year before and after. The format is inspired from [29]

communities.

Previous studies have made notable progress in the implementation of text-to-SQL and have undergone a long evolutionary process. Early efforts were mostly based on well-designed rules and templates [10, 11], specifically suitable for simple database scenarios. In recent years, with the heavy labor costs [12] brought by rule-based methods and the growing complexity of database environments [13–15], designing a rule or template for each scenario has become increasingly difficult and impractical. The development of deep neural networks has advanced the progress of text-to-SQL [16, 17], which can automatically learn a mapping from the user question to its corresponding SQL [18, 19]. Subsequently, pre-trained language models (PLMs) with strong semantic parsing capacity have become the new paradigm for text-to-SQL [20], taking its performance to a new level [21–23]. Incrementally, PLM-based research focused on table content encoding [19, 24, 25] and pre-training [11, 20] has further advanced this field. Recently, the LLM-based approaches implementing text-to-SQL through in-context learning (ICL) [7] and supervised fine-tuning (SFT) [9] paradigm, reaching state-of-the-art accuracy with the well-designed framework and stronger comprehension capability compared to PLMs.

The overall implementation details of LLM-based text-to-SQL can be divided into 3 aspects: **1. Question understanding:** The NL question is a semantic representation of the user’s intention, which the corresponding generated SQL query is expected to align with; **2. Schema comprehension:** The schema provides the table and column structure of the database, and the text-to-SQL system is required to identify the target components in the database that match the user question; **3. SQL generation:** This involves incorporating the above parsing and then writing correct syntax to generate executable SQL queries that can retrieve the desired answer. The LLMs have proven to perform a good vanilla implementation [26, 27], benefiting from the more powerful semantic parsing enabled by the richer training corpus compared to the PLMs [28, 29]. Further studies on enhancing the LLMs for question understanding [7, 8], schema comprehension [30, 31], and SQL generation [32] are being increasingly released.

A. Challenges in Text-to-SQL

Despite the significant progress made in text-to-SQL research, several challenges remain that hinder the development of

robust and generalized text-to-SQL systems [73]. In this survey, we aim to catch up with the recent advances and provide a comprehensive review of the current state-of-the-art (SOTA) models and approaches in LLM-based text-to-SQL. We begin by introducing the fundamental concepts and challenges associated with text-to-SQL, highlighting the importance of this task in various domains. We then delve into the evolution of LLMs and their application to text-to-SQL, discussing the key advancements and breakthroughs in this field. After the overview, we provide a detailed introduction to the recent advances of text-to-SQL incorporating LLMs. Specifically, the body of our survey covers a range of contents related to LLM-based text-to-SQL, including:

- **Datasets and Benchmarks:** We provide an overview of the commonly used datasets and benchmarks for evaluating LLM-based Text-to-SQL systems. We discuss their characteristics, complexity, and the challenges they pose for text-to-SQL system development and evaluation.
- **Evaluation Metrics:** We present the evaluation metrics used to assess the performance of LLM-based Text-to-SQL systems, including accuracy, exactness, and execution correctness. We discuss the advantages and limitations of each metric and their relevance to real-world applications.
- **Methods and Models:** We explore the different methods and models employed for LLM-based text-to-SQL, including in-context learning and fine-tuning-based paradigms. We discuss their implementation details, strengths, and adaptations specific to the text-to-SQL task.
- **Expectations and Future Directions:** We discuss the current challenges and limitations of LLM-based Text-to-SQL, such as real-world robustness, computational efficiency, data privacy and extensions. We also outline potential future research directions and opportunities for improvement.

We hope this survey will provide a clear overview of recent studies and inspire future studies. A taxonomy tree is shown in Fig. 3

II. OVERVIEW

Text-to-SQL is a task that aims to convert natural language-type questions into corresponding SQL queries that can be executed on a relational database. Formally, given a user question Q (also known as a user query, natural language, NL, NL question, etc.) and a database schema \mathcal{D} , the goal

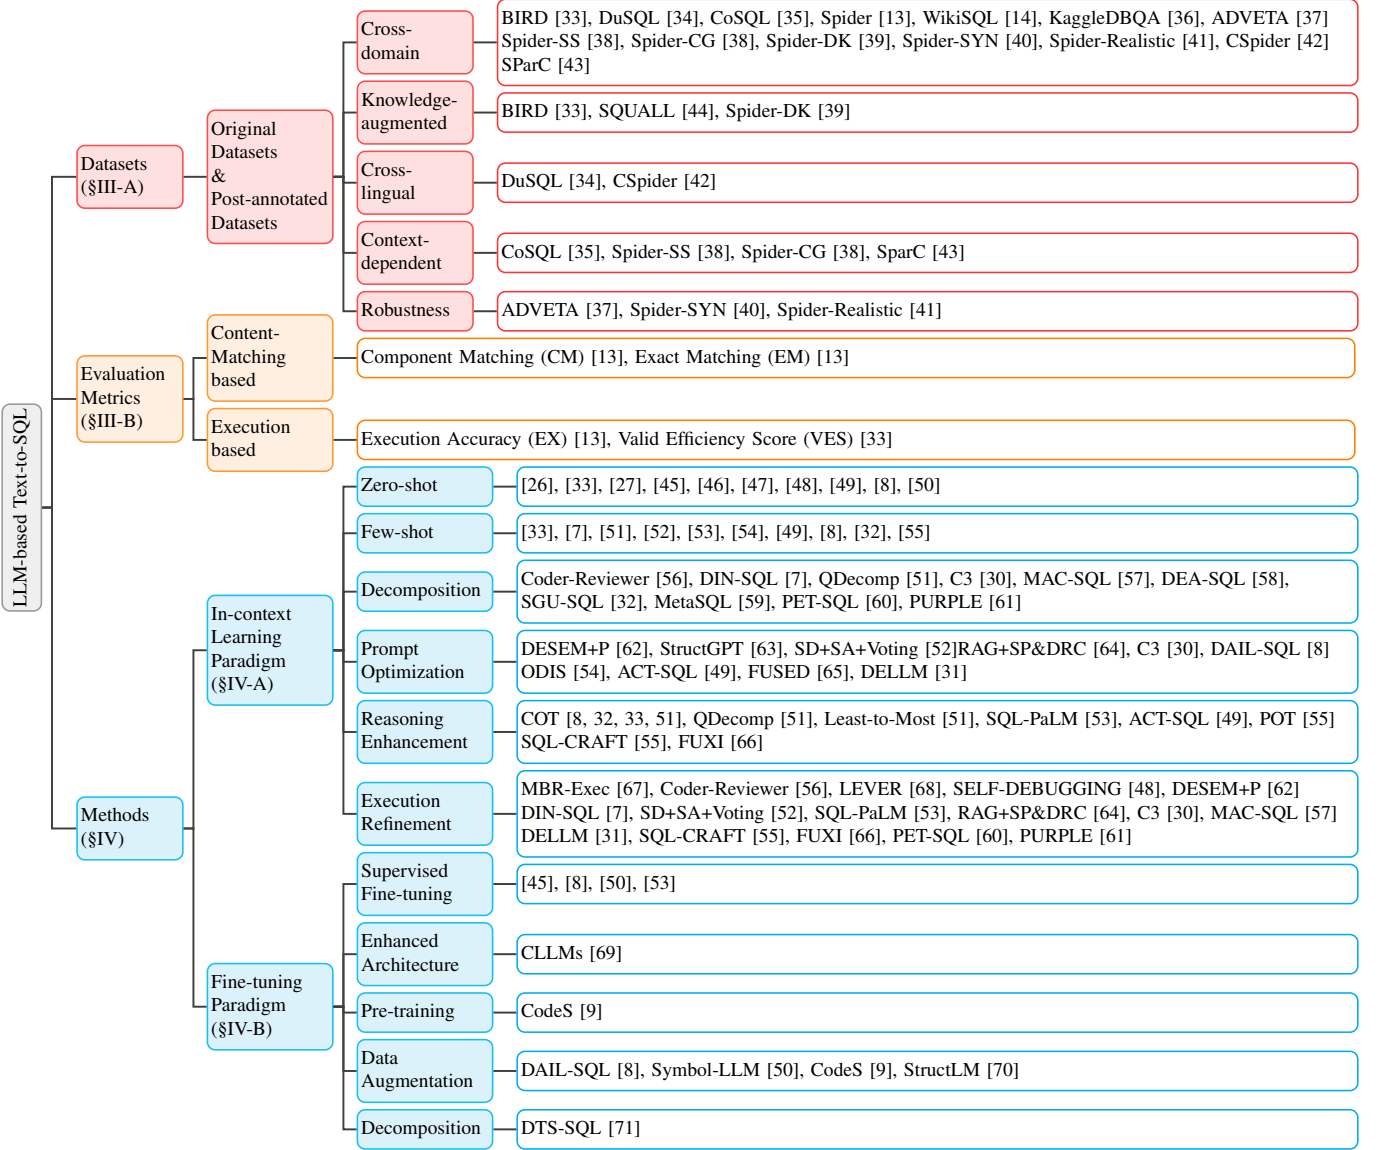


Fig. 3: Taxonomy tree of the research in LLM-based text-to-SQL. The display order in each node is organized by the released time. The format is adapted from [72]

is to generate an SQL query Y that can be executed on the database to obtain the desired answer. Text-to-SQL has the potential to democratize access to data by allowing users to interact with databases using natural language without the need for specialized knowledge of SQL. This can benefit various domains, such as business intelligence, customer support, and scientific research, by enabling non-technical users to easily retrieve information from databases and facilitating more efficient data analysis.

1) Linguistic Complexity and Ambiguity: Natural language questions often contain complex linguistic structures, such as nested clauses, coreferences, and ellipses, which make it challenging to map them accurately to SQL queries. Additionally, natural language is inherently ambiguous, with multiple possible interpretations for a given question. Resolving these ambiguities and understanding the intent behind the question requires deep language understanding and the ability

to incorporate context and domain knowledge.

2) Schema Understanding and Representation: To generate accurate SQL queries, text-to-SQL systems need to have a comprehensive understanding of the database schema, including table names, column names, and relationships between tables. However, database schemas can be complex and vary significantly across different domains. Representing and encoding the schema information in a way that can be effectively utilized by the text-to-SQL model is a challenging task.

3) Rare and Complex SQL Operations: Some SQL queries involve rare or complex operations, such as nested subqueries, outer joins, and window functions. These operations are less frequent in the training data and pose challenges for text-to-SQL models to generate accurately. Designing models that can handle a wide range of SQL operations, including rare and complex ones, is an important consideration.

TABLE I: The statistics and analysis of well-known datasets of text-to-SQL ordered by release time. The original dataset indicates that the dataset is designed with a corresponding database, while post-annotated datasets involve annotating new components within existing datasets and databases rather than releasing a new database.

Original Dataset	Release Time	#Example	#DB	#Table/DB	#Row/DB	Characteristics
BIRD [33]	May-2023	12,751	95	7.3	549K	Cross-domain, Knowledge-augmented
KaggleDBQA [36]	Jun-2021	272	8	2.3	280K	Cross-domain
DuSQL [34]	Nov-2020	23,797	200	4.1	-	Cross-domain, Cross-lingual
SQUALL [44]	Oct-2020	11,468	1,679	1	-	Knowledge-augmented
CoSQL [35]	Sep-2019	15,598	200	-	-	Cross-domain, Context-dependent
Spider [13]	Sep-2018	10,181	200	5.1	2K	Cross-domain
WikiSQL [14]	Aug-2017	80,654	26,521	1	17	Cross-domain
Post-annotated Dataset	Release Time	Base Dataset	Special Setting		Characteristics	
ADVETA [37]	Dec-2022	Spider, etc.	Adversarial table perturbation		Robustness	
Spider-SS&CG [38]	May-2022	Spider	Splitting example into sub-examples		Context-dependent	
Spider-DK [39]	Sep-2021	Spider	Adding domain knowledge		Knowledge-augmented	
Spider-SYN [40]	Jun-2021	Spider	Manual synonym replacement		Robustness	
Spider-Realistic [41]	Oct-2020	Spider	Removing column names in question		Robustness	
CSpider [42]	Sep-2019	Spider	Chinese version of Spider		Cross-lingual	
SPaC [43]	Jun-2019	Spider	Annotate conversational contents		Context-dependent	

4) Cross-Domain Generalization: Text-to-SQL models often struggle to generalize across different database schemas and domains. Models trained on a specific domain may not perform well on questions from a different domain due to differences in vocabulary, schema structure, and question patterns. Developing models that can effectively adapt to new domains with minimal fine-tuning or domain-specific training data is an ongoing challenge.

A. Evolutionary Process

The field of text-to-SQL has witnessed significant advancements over the years, evolving from rule-based approaches to deep learning-based methods and, more recently, to the integration of pre-trained language models (PLMs) and large language models (LLMs), as shown in Fig. 2.

1) Rule-based Approaches: Early text-to-SQL systems relied heavily on rule-based approaches [10–12], where manually crafted rules and heuristics were used to map natural language questions to SQL queries. These approaches often involved extensive feature engineering and domain-specific knowledge. While rule-based methods achieved success in specific domains, they lacked the flexibility and generalization capabilities needed to handle diverse and complex questions.

2) Deep Learning-based Methods: With the rise of deep learning, sequence-to-sequence models, such as LSTMs and transformers, were adapted to generate SQL queries from natural language input [19, 74]. Typically, RYANSQL [19] introduced techniques like intermediate representations and sketch-based slot filling to handle complex questions and improve cross-domain generalization. Recently, researchers introduced graph neural networks (GNNs) for text-to-SQL tasks by leveraging schema dependency graphs to capture the relationships between database elements [18, 75].

3) PLM-based Approaches: Pre-trained language models (PLMs) have emerged as a powerful solution for text-to-SQL, leveraging the vast amounts of linguistic knowledge and understanding captured during pre-training. The early adoption of PLMs in text-to-SQL primarily focused on fine-tuning pre-

existing PLMs, such as BERT [24] and RoBERTa [76], on text-to-SQL datasets. These PLMs, which were pre-trained on large amounts of text data, captured rich semantic representations and language understanding capabilities. By fine-tuning them on text-to-SQL tasks, researchers aimed to leverage the knowledge and linguistic understanding of PLMs to generate accurate SQL queries [20, 74]. Another line of research focuses on incorporating schema information into PLMs to improve their understanding of database structures and generate more accurate SQL queries. Schema-aware PLMs are designed to capture the relationships and constraints present in the database schema [21].

4) LLM-based Implementation: Large language models (LLMs), such as GPT series [77, 78], have gained significant attention in recent years due to their ability to generate coherent and fluent text. Researchers have started exploring the potential of LLMs for text-to-SQL by leveraging their vast knowledge and generative capabilities [8, 26]. These approaches often involve prompt engineering to guide the proprietary LLMs in the generation process or fine-tuning the open-source LLMs on text-to-SQL datasets.

The integration of LLMs in text-to-SQL is still an emerging area, and there is significant potential for further exploration and improvement. Researchers are investigating ways to better leverage the knowledge and reasoning capabilities of LLMs, incorporate domain-specific knowledge, and develop more efficient fine-tuning strategies. As the field continues to evolve, we can expect to see more advanced and powerful LLM-based approaches that push the boundaries of text-to-SQL performance and generalization.

III. BENCHMARKS & EVALUATION

In this section, we introduce the benchmarks for text-to-SQL, encompassing well-known datasets and evaluation metrics.

A. Datasets

As shown in Table I, we classify the datasets into "Original Datasets" and "Post annotated Datasets" based on whether

they were released with the original dataset and databases or created by adapting existing datasets and databases with special processing. We also highlight their release times. For the original datasets, we provide a detailed analysis, including the number of examples, databases, tables per database, and rows per database. For the post-annotated datasets, we identify their base dataset and describe the special processing applied to them. To illustrate the potential opportunities of each dataset, we categorize their characteristics into Cross-domain, Knowledge-augmented, Context-dependent, Robustness, and Cross-lingual, which we will discuss in detail below.

1) Cross-domain Dataset. This refers to datasets where the background information of different databases comes from various domains. Since real-world text-to-SQL tasks often involve databases from multiple domains, most original text-to-SQL datasets [13, 14, 33–36] and post annotated datasets [37–43] are in the cross-domain setting to fit well with the needs of cross-domain applications.

2) Knowledge-augmented Dataset. Interest in incorporating knowledge into text-to-SQL tasks has increased significantly in recent years. BIRD [33] employs domain experts to annotate each text-to-SQL sample with external knowledge, categorized into Numeric Reasoning Knowledge, Domain Knowledge, Synonym Knowledge, and Value Illustration. Similarly, Spider-DK [39] defines and adds five types of domain knowledge for a human-curated version of the Spider dataset: SELECT Columns Mentioned by Omission, Simple Inference Required, Synonyms Substitution in Cell Value Word, One Non-Cell Value Word Generate a Condition, and Easy to Conflict with Other Domains. Both studies found that human-annotated knowledge significantly improves SQL generation accuracy for samples requiring external domain knowledge. Additionally, SQUALL [44] manually annotates alignments between the words in NL questions and the entities in SQL, providing finer-grained supervision than other datasets.

3) Context-dependent Dataset. SPaRC [43] and CoSQL [35] explore context-dependent SQL generation by constructing a conversational database querying system. Unlike traditional text-to-SQL datasets that only have a single (question, SQL) pair for one example, SPaRC decomposes the (question, SQL) examples in the Spider dataset into multiple (sub-question, SQL) pairs to construct a simulated and meaningful interaction, including inter-related sub-questions that aid SQL generation, and unrelated sub-questions that enhance data diversity. CoSQL, in comparison, involves conversational interactions in natural language, simulating real-life scenarios to increase complexity and diversity. Additionally, Spider-SS&CG [38] splits the NL question in the Spider dataset into multiple sub-questions and sub-SQLs, demonstrating that training on these sub-examples can improve a text-to-SQL model’s generalization ability on out-of-distribution samples.

4) Robustness Dataset. Evaluating the accuracy of text-to-SQL models with attacked or perturbed database contents (e.g., schema and tables) is crucial for assessing robustness. Spider-Realistic [41] removes explicit schema-related words from the NL questions, while Spider-SYN [40] replaces them with manually selected synonyms. ADVETA [37] introduces Adversarial Table Perturbation (ATP), which perturbs tables by

replacing original column names with misleading alternatives and inserting new columns with high semantic associations but low semantic equivalency. These perturbations lead to significant drops in accuracy, as a text-to-SQL model with low robustness may be misled by incorrect matches between tokens in NL questions and database entities.

5) Cross-lingual Dataset. SQL keywords, function names, table names, and column names are typically written in English, posing challenges for applications in other languages. CS Spider [42] translates the Spider dataset into Chinese, identifying new challenges in word segmentation and cross-lingual matching between Chinese questions and English database contents. DuSQL [34] introduces a practical text-to-SQL dataset with Chinese questions and database contents provided in both English and Chinese.

B. Evaluation Metrics

We introduce four widely used evaluation metrics for the text-to-SQL task as follows: Component Matching and Exact Matching, which are based on SQL content matching, and Execution Accuracy and Valid Efficiency Score, which are based on execution results.

1) Content Matching-based Metrics. SQL content matching metrics focus on comparing the generated SQL query with the ground truth SQL query based on their structural and syntactic similarities.

- **Component Matching (CM)** [13] evaluates the model performance by measuring the exact match between predicted and ground truth SQL components—SELECT, WHERE, GROUP BY, ORDER BY, and KEYWORDS—using the F1 score. Each component is decomposed into sets of sub-components and compared for an exact match, accounting for SQL components without order constraints.
- **Exact Matching (EM)** [13] measures the percentage of examples whose predicted SQL query is equivalent to the ground truth SQL query. A predicted SQL is considered correct only if all its components, as described in Component Matching, match exactly with those of the ground truth query.

2) Execution-based Metrics. Execution result metrics assess the correctness of the generated SQL query by comparing the results obtained from executing the query on the target database with the expected results.

- **Execution Accuracy (EX)** [13] measures the correctness of a predicted SQL query by executing it in the corresponding database and comparing the executed results with the ground truth queries.
- **Valid Efficiency Score (VES)** [33] is defined to measure the efficiency of valid SQL queries, which are the predicted SQL queries whose executed results exactly match the ground truth query. Thus, VES evaluates both the efficiency and accuracy of predicted SQL queries. For an evaluation dataset with N examples, VES can be computed by:

$$\text{VES} = \frac{1}{N} \sum_{n=1}^N \mathbb{1}(V_n, \hat{V}_n) \cdot \mathbf{R}(Y_n, \hat{Y}_n), \quad (1)$$

where \hat{Y}_n and \hat{V}_n are the predicted SQL query and its executed results. Y_n and V_n are those of the corresponding ground truth SQL query. $\mathbb{1}(V_n, \hat{V}_n)$ is an indicator function, where:

$$\mathbb{1}(V_n, \hat{V}_n) = \begin{cases} 1, & V_n = \hat{V}_n \\ 0, & V_n \neq \hat{V}_n \end{cases} \quad (2)$$

Then, $\mathbf{R}(Y_n, \hat{Y}_n) = \sqrt{E(Y_n)/E(\hat{Y}_n)}$ denotes the relative execution efficiency of the predicted SQL query in comparison to ground-truth query, where $E(\cdot)$ is the execution time of each SQL in the database. [33] ensures the stability of this metric by computing the average of $\mathbf{R}(Y_n, \hat{Y}_n)$ over 100 runs for each example.

Most of the recent LLM-based text-to-SQL studies focus on these four datasets: Spider [13], Spider-Realistic [41], Spider-SYN [40], and BIRD [33]; and these three evaluation methods: EM, EX, and VES, we will primarily focus on them in the following introduction.

IV. METHODS

The implementation of LLM-based applications mostly relies on in-context learning (prompt engineering) [79] and fine-tuning [80] since the powerful proprietary models and well-architected open-source models are being released in large quantities [81–83]. The specific methods can generally be divided into two paradigms, in-context learning (ICL) and fine-tuning (FT). In this survey, we primarily focus on these text-to-SQL paradigms and will discuss them accordingly.

A. In-context Learning

Through extensive and widely recognized research, prompt engineering has been proven to play a decisive role in the performance of LLMs [28], also impacting the text-to-SQL task in different prompt styles [8, 46]. Necessarily, developing text-to-SQL methods in the in-context learning (ICL) paradigm is valuable for achieving promising improvement. The implementation of LLM-based text-to-SQL process to generate executable SQL query Y is formulated as:

$$Y = f(Q, S, \mathcal{I} \mid \theta), \quad (3)$$

where Q represents the user question. S is the database schema/content, which can be decomposed as $S = \langle \mathcal{C}, \mathcal{T}, \mathcal{K} \rangle$, where $\mathcal{C} = \{c_1, c_2, \dots\}$ and tables $\mathcal{T} = \{t_1, t_2, \dots\}$ represent the collection of various columns and tables, \mathcal{K} is the potentially external knowledge (e.g. foreign key relationships [49], schema linking [30] and domain knowledge [31, 33]). \mathcal{I} represents the instruction for the text-to-SQL task, which performs indicative guidance to guide the LLMs for generating an accurate SQL query. $f(\cdot \mid \theta)$ is a LLM with parameter θ . In the in-context learning (ICL) paradigm, we utilize an off-the-shelf text-to-SQL model (i.e., parameter θ of the model is frozen) for implementation. Various well-designed methods have been adopted in the ICL paradigm for text-to-SQL. We group them into five categories $\mathbf{C}_{0:4}$, including: \mathbf{C}_0 -Trivial Prompt, \mathbf{C}_1 -Decomposition, \mathbf{C}_2 -Prompt Optimization, \mathbf{C}_3 -Reasoning Enhancement, and \mathbf{C}_4 -Execution Refinement, the details are shown in Tab. II.

TABLE II: Typical methods used for in-context learning (ICL) in LLM-based text-to-SQL. The full table of existing methods with categorization $\mathbf{C}_{1:4}$ and more details are listed in Table III.

Methods	Adopted by	Applied LLMs
\mathbf{C}_0 -Trivial Prompt	Zero-shot [26] Few-shot [8]	ChatGPT ChatGPT
\mathbf{C}_1 -Decomposition	DIN-SQL [7]	GPT-4
\mathbf{C}_2 -Prompt Optimization	DAIL-SQL [8]	GPT-4
\mathbf{C}_3 -Reasoning Enhancement	ACT-SQL [49]	GPT-4
\mathbf{C}_4 -Execution Refinement	LEVER [68]	Codex

\mathbf{C}_0 -Trivial Prompt: Trained through massive data, LLMs have a strong overall proficiency in different downstream tasks with zero-shot and few-shot prompting [80, 84, 85], which is widely recognized and achieves promising results. In our survey, we categorized the above prompting approaches without the well-designed framework as trivial prompts. As introduced above, the formulated process of LLM-based text-to-SQL Eq. 3 can also represent zero-shot prompting, the input \mathcal{P}_0 can be obtained by concatenate \mathcal{I} , S and Q :

$$\mathcal{P}_0 = \mathcal{I} \oplus S \oplus Q. \quad (4)$$

To regulate the prompting process, we set the OpenAI demonstration² as the standard (trivial) prompt [30] for text-to-SQL.

Zero-shot: Many research works [26, 27, 46] utilize zero-shot prompting, focusing mainly on the influence of the style of prompt construction and the zero-shot performance of various LLMs for text-to-SQL. As an empirical evaluation, [26] evaluates the baseline text-to-SQL capabilities of different early-developed LLMs [77, 86, 87], which also presents the results for different prompt styles. The results indicate that prompt design is critical for performance, with error analysis, [26] propose more database content can harm the overall performance. Since ChatGPT emerged with impressive capabilities in conversational scenarios and code generation, [27] assesses its performance of text-to-SQL. With zero-shot settings, the results demonstrate that ChatGPT has a promising text-to-SQL ability compared to the state-of-the-art (SOTA) models. For fair comparability, [47] reveal effective prompt construction for LLM-based text-to-SQL; they study different styles of prompt construction and make conclusions of zero-shot prompt design based on the comparisons.

Primary key and foreign key carry contiguous knowledge of different tables, [49] study their impact by adding these keys on different designed database prompt style on zero-shot prompting results. A benchmark evaluation [8] also studies the influence of foreign keys, with five different prompt representation styles, each style can be considered as the permutation and combination of the instruction, rule implication, and foreign key. Apart from the foreign key, this study also explores zero-shot prompting combined with “no explanation” rule and the rule implication “Let’s think step by step”. Empowered by the annotated knowledge of human experts, [33] follow the standard prompting and obtain improvement by combining the provided oracle knowledge. With the explosion of open-source LLMs,

²The prompt style that follows the official document from OpenAI platform: <https://platform.openai.com/examples/default-sql-translate>

TABLE III: Well-designed methods used in in-context learning (ICL) paradigm for LLM-based text-to-SQL ordered by release time. The methods are grouped in four categories based on their implementation perspective: C_1 -Decomposition, C_2 -Context Augmentation, C_3 -Reasoning Enhancement, C_4 -Execution Refinement. The method in multiple categories will be introduced respectively. *There are multiple applied LLMs in the corresponding method; we present the selected LLM with representative performance. †COT method are reported in multiple venues: NeurIPS’23 [33], EMNLP’23 [51], VLDB’24 [8], arXiv’24 [32]

Methods	Applied LLMs	Benchmark	Metrics	C_1	C_2	C_3	C_4	Release Time	Publication Venue
MBR-Exec [67]	Codex	[13]	EX				✓	Apr-2022	EMNLP’22
Coder-Reviewer [56]	Codex	[13]	EX	✓			✓	Nov-2022	ICML’23
LEVER [68]	Codex	[13]	EX				✓	Feb-2023	ICML’23
SELF-DEBUGGING [48]	StarCoder*	[13]	EX				✓	Apr-2023	ICLR’24
DESEM+P [62]	ChatGPT	[13, 40]	EX		✓		✓	Apr-2023	PRICAI’23
DIN-SQL [7]	GPT-4*	[13, 33]	EX, EM, VES	✓			✓	Apr-2023	NeurIPS’23
COT [8, 32, 33, 51]	GPT-4	[13, 33, 41]	EX, VES			✓		May-2023	Multiple Venues†
StructGPT [63]	ChatGPT*	[13, 40, 41]	EX		✓			May-2023	EMNLP’23
SD+SA+Voting [52]	ChatGPT*	[13, 40, 41]	EX		✓		✓	May-2023	EMNLP’23 Findings
QDecomp [51]	Codex	[13, 41]	EX	✓		✓		May-2023	EMNLP’23
Least-to-Most [51]	Codex	[13]	EX			✓		May-2023	EMNLP’23
SQL-PaLM [53]	PaLM-2	[13]	EX			✓	✓	May-2023	arXiv’23
RAG+SP&DRC [64]	ChatGPT	[13]	EX		✓		✓	Jul-2023	ICONIP’23
C3 [30]	ChatGPT	[13]	EX	✓	✓		✓	Jul-2023	arXiv’23
DAIL-SQL [8]	GPT-4*	[13, 33, 41]	EX, EM, VES		✓			Aug-2023	VLDB’24
ODIS [54]	Codex*	[13]	EX		✓			Oct-2023	EMNLP’23 Findings
ACT-SQL [49]	GPT-4*	[13, 40]	EX, EM		✓	✓		Oct-2023	EMNLP’23 Findings
MAC-SQL [57]	GPT-4*	[13, 33]	EX, EM, VES	✓			✓	Dec-2023	arXiv’23
DEA-SQL [58]	GPT-4	[13]	EX	✓				Feb-2024	ACL’24 Findings
FUSED [65]	ChatGPT*	[13]	EX		✓			Feb-2024	arXiv’24
DELLM [31]	GPT-4*	[13, 33]	EX, VES		✓		✓	Feb-2024	ACL’24 Findings
SGU-SQL [32]	GPT-4*	[13, 33]	EX, EM	✓				Feb-2024	arXiv’24
POT [55]	GPT-4*	[13, 33]	EX			✓		Feb-2024	arXiv’24
SQL-CRAFT [55]	GPT-4*	[13, 33]	EX			✓	✓	Feb-2024	arXiv’24
FUXI [66]	GPT-4*	[33]	EX			✓	✓	Feb-2024	arXiv’24
MetaSQL [59]	GPT-4*	[13]	EX, EM	✓				Feb-2024	ICDE’24
PET-SQL [60]	GPT-4	[13]	EX	✓			✓	Mar-2024	arXiv’24
PURPLE [61]	GPT-4*	[13, 40, 41]	EX, EM	✓			✓	Mar-2024	ICDE’24

according to the results of similar evaluation, these models are also capable of zero-shot text-to-SQL task [45, 46, 50], especially code generation models [46, 48]. For zero-shot prompting optimization, [46] raises a challenge for designing an effective prompt template for LLMs; the former prompt construction lacks structure uniformity, which makes it hard to find out a concrete element within a prompt constructing template influences the performance of LLMs. They address this challenge by investigating a more unified series of prompt templates warping with different prefixes, infixes, and postfixes. **Few-shot:** The technique of few-shot prompting is widely used in both practical applications and well-designed research, which has been proven efficient for eliciting better performance of LLMs [28, 88]. The input prompt of the few-shot approach LLM-based text-to-SQL can be formulated as an extension of Eq. 3:

$$\mathcal{P}_n = \{\mathcal{F}_1, \mathcal{F}_2, \dots, \mathcal{F}_n\} \oplus \mathcal{P}_0, \quad (5)$$

where the \mathcal{P}_n represent the input prompt for n -shot learning, n is the provided instances (examples) number; \mathcal{F} denote the few-shot instance, which can be decomposed as $\mathcal{F}_i = (\mathcal{S}_i, \mathcal{Q}_i, \mathcal{Y}_i)$, i is the serial number of instances. The study of few-shot prompting focuses on the number of representations and few-shot instance selection.

As pilot experiments, few-shot prompting for text-to-SQL are evaluated in multiple datasets with various LLMs [7, 32], achieve better performance rather than zero-shot approaches. [33] provides a 1-shot example combining with CoT [89]

prompting, the decomposed reasoning steps of the given example trigger the text-to-SQL model for generating accurate SQL. [55] study the effect of the number of few-shot examples. [52] focus on the sampling strategies by studying the similarity and the diversity between different demonstrations, setting random sampling as the baseline, and evaluating different strategies and their combination for comparison. Furthermore, above the similarity selection, [8] evaluated masked question similarity selection and the upper limit of similarity approaches with various numbers of few-shot examples. A study of difficult-level samples selection [51] compared the performance of few-shot Codex, with random selection and difficulty-based selection for few-shot instances on difficulty categorized dataset [13, 41]. Three difficulty-based selection strategies are devised based on the number of selected samples at different difficulty levels. [49] utilize a hybrid strategy for selecting samples, which combines static examples and similarity-based dynamic examples for few-shot prompting. In their evaluations, they also test the impact of different input schema styles and different static and dynamic exemplar numbers.

The impact of cross-domain few-shot examples is also being studied [54]. When incorporating in-domain and out-of-domain with different numbers of examples, the in-domain demonstration always outperforms zero-shot and out-of-domain examples and gets better as the number of examples rises. To explore the detailed construction of input prompt, [53] compare the concise and verbose prompt design approaches. The former

style splits the schema, the column names, and the primary and foreign keys by bar, and the latter organizes them as natural language descriptions.

C₁-Decomposition: As an intuitive solution, decomposing a challenging user question into simpler sub-questions and using multi-step reasoning for implementation can reduce the complexity of the full text-to-SQL task [7, 51]. Dealing with less complexity, LLMs have the potential to have better performance. The decomposition approaches for LLM-based text-to-SQL are categorized into two paradigms: (1) **sub-task decomposing**, provides additional parsing to assist the final SQL generation by decomposing the overall text-to-SQL task into smaller effective sub-tasks (e.g., schema linking, domain classification). (2) **sub-question decomposing**, divides the user question into sub-questions to reduce the question’s complexity and difficulty, then generates the sub-SQL by solving these questions to deduce the final SQL query.

DIN-SQL [7] proposed a decomposed in-context learning method consisting of four modules: schema linking, classification & decomposition, SQL generation, and self-correction. DIN-SQL first finishes the schema linking between the user question and the target database; the following module decomposes the user question into correlated sub-questions and does a difficulty classification. Based on the above information, the SQL generation module generates a corresponding SQL, and the self-correction module identifies and corrects the potential errors in the predicted SQL. This approach comprehensively considers the decomposition of both sub-tasks and sub-questions. Coder-Reviewer [56] framework proposed a re-ranking method, combining Coder models for the generation and Reviewer models to evaluate the likelihood of the instruction. Refer to the Chain-of-Thought [89] and Least-to-Most prompting [90], QDecomp [51] introduce question decomposition prompting, which follows the question reduction stage in least-to-most prompting and instruct the LLM to decompose the original complex question as the intermediate reasoning steps. C3 [30] consists of three key components: clear prompting, calibration bias prompting, and consistency; these components are accomplished by assigning ChatGPT with different tasks. Firstly, the clear prompting component generates the schema linking and the distilled question-relevant schema as a clear prompt. Then, a multi-turn dialogue about text-to-SQL hints is utilized as a calibration bias prompt, which combines with the clear prompt to guide the SQL generation. The generated SQL queries are selected by consistency and execution-based voting to get the final SQL. MAC-SQL [57] presents a multi-agent collaborating framework; the text-to-SQL process is finished as the collaboration of the agents: Selector, Decomposer, and Refiner. The Selector preserves relevant tables for user questions; the Decomposer breaks down user questions into sub-questions and provides solutions; finally, the Refiner validates and refines the defective SQL. DEA-SQL [58] introduces a workflow paradigm aiming to enhance the attention and problem-solving scope of LLM-based text-to-SQL through decomposition. This method decomposes the overall task, enabling the SQL generation module to have the corresponding prerequisite (information determination, question classification) and subsequent (self-correction, active learning)

sub-tasks. Through the workflow paradigm, an accurate SQL query is generated. SGU-SQL [32] is a structure-to-SQL framework, leveraging the inherent structure information to assist SQL generation. Specifically, the framework constructs a graph structure for the user question and the corresponding database respectively, then uses the encoded graphs to construct structure linking [91, 92]. A meta operator decomposes the user question with a grammar tree and finally designs the input prompt with meta-operation in SQL. MetaSQL [59] introduces a three-stage approach for SQL generation, which consists of decomposition, generation, and rank. The decomposition stage uses semantic decomposition and metadata composition to process the user question. Taking the previously processed data as input, a text-to-SQL model using metadata-conditioned generation to generate some candidate SQL queries. Finally, a two-stage ranking pipeline is applied to get a global-optimal SQL query. PET-SQL [60] proposed a prompt-enhanced two-stage framework. Firstly, an elaborated prompt instructs the LLMs to generate preliminary SQL (PreSQL) where some few-shot demonstrations are selected based on similarity. Then, schema linking is found based on PreSQL and combined to prompt the LLMs to generate the Final SQL (FinSQL). Finally, multiple LLMs are utilized to generate a FinSQL, ensuring consistency based on the execution results.

C₂-Prompt Optimization: As previously introduced, few-shot learning is widely studied for prompting LLMs [77]. For LLM-based text-to-SQL with in-context learning, trivial few-shot approaches obtained promising results [7, 8, 33], further optimization of few-shot prompting has the potential to achieve higher performance. Since the performance of SQL generation in off-the-shelf LLMs largely depends on the quality of the corresponding input prompt [93], many decisive factors that can influence the quality of the prompt have become focuses of the research [8] (e.g., quality and quantity in the few-shot organization, the similarity between user questions and few-shots instances, external knowledge/hints). The process of prompt quality improvement is actually the prompt’s optimization, including **few-shot sampling strategies**, **schema augmentation**, and **external knowledge generation**.

DESEM [62] is a prompt engineering framework with de-semanticization and skeleton retrieval. The framework first employs domain-specific words masking module to remove the semantic tokens in questions that preserve the question’s intentions. And then utilizes an adjustable prompting module that retrieves the few-shot examples with identical question intentions and incorporates schema-relevance filtering to guide the LLM’s SQL generation. The QDecomp [51] framework introduces the InterCOL mechanism to incrementally incorporate the decomposed sub-questions with correlative table and column names. With difficulty-based selection, the few-shot examples for QDecomp are difficult-level sampled. Besides similarity-diversity sampling, [52] proposed SD+SA+Voting (Similarity-Diversity+Schema Augmentation+Voting) sampling strategy. They first employ semantic similarity and k -Means cluster diversity for sampling few-shot examples and then enhance the prompt with schema knowledge (semantic or structure augmentation). C3 [30] framework comprises a clear prompting component, which takes the question and schema as

the LLMs input, generates a clear prompt that includes a schema that removes the redundant information irrelevant to the user question and a schema linking, and also a calibration component providing hints. The LLMs take their composition as context-augmented prompts for SQL generation. A retrieval-augmented framework is introduced with sample-aware prompting [64], which simplifies the original question and extracts the question skeleton from the simplified question, then finishes the sample retrieval in the repository according to skeleton similarities. The retrieved samples are combined with the original question for few-shot prompting. ODIS [54] introduces the selection of a sample with out-of-domain demonstrations and in-domain synthetic data, which retrieves few-shot demonstrations from hybrid sources to augment the prompt representations. DAIL-SQL [8] proposed a novel approach to address the issues in few-shot sampling and organization, presenting a better balance between the quality and quantity of few-shot examples. DAIL Selection first masks domain-specific words in user and few-shot example questions, then ranks the candidate examples based on the embedded Euclidean distance. Meanwhile, the similarity between the pre-predicted SQL queries is calculated. Finally, the selection mechanism obtains the similarity-sorted candidates according to the pre-set criteria. The few-shot examples are guaranteed good similarity with both questions and SQL queries with this method. ACT-SQL [49] proposed dynamic examples in few-shot prompting, which is selected according to similarity score. FUSED [65] are presented to build a high-diversity demonstrations pool through human-free multiple-iteration synthesis to improve the diversity of the few-shot demonstrations. The pipeline of FUSED samples the demonstrations to be fused by clustering, then fuse the sampled demonstrations to construct the pool to enhance few-shot learning. Knowledge-to-SQL [31] framework aims to build a Data Expert LLM (DELLM) to provide knowledge for SQL generation. The DELLM is trained by supervised fine-tuning using human expert annotations [33] and further refined by preference learning with the database’s feedback. DELLM generates four categories of knowledge, the well-designed methods (e.g. DAIL-SQL [8], MAC-SQL [57]) incorporating the generated knowledge to achieve better performance for LLM-based text-to-SQL with in-context learning.

C₃-Reasoning Enhancement: LLMs have exhibited promising capabilities in tasks involving commonsense reasoning, symbolic reasoning, and arithmetic reasoning [94]. Since for the text-to-SQL tasks, numeric and synonym reasoning frequently occur in realistic scenarios [33, 41], the prompting strategies for the LLMs reasoning possess the potential to enhance their SQL generation capabilities. Recent studies primarily focus on incorporating well-designed reasoning-enhanced methods for text-to-SQL adaptation, improving LLMs to address the challenge about **complex questions that require multi-step reasoning** and the issue of **self-consistency** [95] in SQL generation.

Chain-of-Thoughts (CoT) prompting technique [89] involves a comprehensive reasoning process that guides LLMs towards accurate deduction, eliciting reasoning in LLMs. The study of LLM-based text-to-SQL utilizes CoT prompting as rule implication [8], which setting the instruction “*Let’s*

think step by step” in prompt construction [8, 32, 33, 51]. However, the straightforward (original) CoT strategy has not demonstrated the potential in text-to-SQL tasks that it has in other reasoning tasks; studying CoT for adaptations is still an ongoing research [51]. Since CoT prompting always uses static examples with human annotation for demonstrations, which requires empirical judgment for the effective selection of few-shot examples, and manual annotating is also an essential need. As a solution, ACT-SQL [49] proposed a method to generate CoT examples automatically. Specifically, given a question, ACT-SQL truncates a set of slices of the question and then enumerates every column appearing in the corresponding SQL query. Each column will be linked with its most relevant slice through the similarity function and appended to the CoT prompt. Through systematical study for enhancing LLMs SQL generation incorporating CoT prompting, QDecomp [51] presents a novel framework to address the challenge for CoT how to come up with the reasoning steps to predict the SQL query. The framework utilizes every slice of the SQL query to construct a logical step in CoT reasoning, then employs natural language templates to articulate each slice of the SQL query and arranges them in the logical execution order. **Least-to-Most** [90] is another prompting technique that decomposes questions into sub-questions and then sequentially solves them. As iterative prompting, pilot experiments [51] demonstrate that it may be unnecessary for text-to-SQL parsing. Using detailed reasoning steps tends to have more error propagation issues. As a variant of CoT, **Program-of-Thoughts (PoT)** prompting strategy [96] are proposed to enhance arithmetic reasoning for LLMs. Through evaluation [55], PoT enhances the LLM for SQL generation, especially in complicated datasets [33]. SQL-CRAFT [55] are proposed to enhance LLM-based SQL generation, which incorporates PoT prompting for Python-enhanced reasoning. PoT strategy requires the model to simultaneously generate the Python code and the SQL queries, enforcing the model to incorporate Python code in its reasoning process. **Self-Consistency** [95] is a prompting strategy improving reasoning in LLMs, which leverages the intuition that a complex reasoning problem typically admits multiple different ways of thinking, leading to its unique correct answer. In the text-to-SQL task, self-consistency is adapted to sampling a set of different SQL and voting for consistent SQL via execution feedback [30, 53]. Similarly, the SD+SA+Voting [52] framework eliminates those with execution errors identified by the deterministic database management system (DBMS) and opts for the prediction that garners the majority vote. Furthermore, motivated by recent research on extending the capabilities of LLMs with tools, FUXI [66] are proposed to enhance LLMs SQL generation through effectively invoking crafted tools.

C₄-Execution Refinement: To design criteria for accurate SQL generation, whether a generated SQL can be successfully executed and elicit a correct answer for the user question is always the priority [13]. As a complex programming task, generating the correct SQL in one go becomes challenging. Intuitively, considering the execution feedback/results in SQL generation assists the alignment to the corresponding database environment, which allows the LLMs to gather the potential

executed errors and results to refine the generated SQL or hold a majority vote [30]. The execution-aware methods in text-to-SQL incorporate the execution feedback in two main approaches: **1) Incorporating the feedback through second round prompting for regeneration**, for every SQL query generated in the initial response, it will be executed in the corresponding database, thus obtaining feedback from the database. This feedback might be an error, or it might yield results that will be appended to the second round prompt. Through in-context learning of this feedback, LLMs are able to refine or regenerate the original SQL, thereby enhancing accuracy. **2) Utilize execution-based selection strategies for generated SQL**, sample multiple generated SQL queries from LLM), and execute each in the database. Based on the results of each SQL query, use selection strategies (e.g., self-consistency, majority vote [60]) to define a query from the SQL set that satisfies the criteria as the final predicted SQL.

MRC-EXEC [67] introduced a natural language to code (NL2Code) translation framework with execution, which executes each sampled SQL query and selects the example with the minimal execution result-based Bayes risk [97]. LEVER [68] proposed an approach to verify NL2Code with execution, utilizing a generation and execution module to collect sampled SQL set and their execution results, respectively, then using a learned verifier to output the probability of the correctness. Similarly, the SELF-DEBUGGING [48] framework is presented to teach LLMs to debug their predicted SQL via few-shot demonstrations. The model is able to refine its mistakes by investigating the execution results and explaining the generated SQL in natural language without human interventions

As previously introduced, to incorporate the well-designed framework with execution feedback, two-stage implications are widely-used: 1. sampling a set of SQL queries. 2. majority vote (self-consistency). Specifically, the C3 [30] framework removes the errors and identifies the most consistent SQL; The retrieval-augmented framework [64] introduced a dynamic revision chain, combining fine-grained execution messages with database content to prompt the LLMs to convert the generated SQL query into natural language explanation; the LLMs are requested to identify the semantic gaps and revise their own generated SQL. Although schema-filtering approaches enhance SQL generation, the generated SQL could be unexecutable. DESEM [62] incorporates a fallback revision to address the issue; it revises and regenerates the SQL base on different kinds of errors and sets termination criteria to avoid the loop. DIN-SQL [7] designed a generic and gentle prompt in their self-correction module; the generic prompt requests the LLMs to identify and correct the error, and the gentle prompt asks the model to check the potential issue. The multi-agent framework MAC-SQL [57] comprises a refiner agent, which is able to detect and automatically rectify SQL errors, taking SQLite error and exception class to regenerate fixed SQL. Since different questions may require different numbers of revisions, SQL-CRAFT [55] framework introduced interactive correction with an automated control determination process to avoid over-correction or insufficient correction. FUXI [66] considers the error feedback in tool-based reasoning for SQL generation. The Knowledge-to-SQL [31] introduces a preference learning

framework incorporating the database execution feedback with a direct preference optimization [98] for refining the proposed DELLM. PET-SQL [60] proposed cross consistency, which comprises two variants: 1) naive voting: instruct multiple LLMs to generate the SQL query, then utilizing the majority vote for the final SQL base on different execution results; 2) fine-grained voting: refine the naive voting based on the difficulty level to mitigate the voting bias.

B. Fine-tuning

Since **supervised fine-tuning** (SFT) is the mainstream approach in the LLMs training [29], for open-source LLMs (e.g., LLaMA-2 [82], Gemma [99]), the most straightforward method to enable the model to adapt a specific domain quickly is to use collected domain label to perform SFT on the model. The SFT phase is typically the preliminary phase of the well-designed training framework [98, 100], as well as the fine-tuning of text-to-SQL. The auto-regressive generation process of SQL query Y can be formulated as follows:

$$P_{\pi}(Y | \mathcal{P}) = \prod_{k=1}^n P_{\pi}(y_k | \mathcal{P}, Y_{1:k-1}), \quad (6)$$

where $Y = \{y_1, y_2, \dots, y_n\}$ is an SQL query of length n , y_k is the corresponding k^{th} token of the SQL query, $Y_{1:k-1}$ is the prefix sequence of Y ahead the token y_k . $P_{\pi}(y_k | \cdot)$ is a conditional probability of a LLM π for generating the k^{th} token of Y base on the input prompt \mathcal{P} and the prefix sequence.

Given a basic open-source model π^0 , the goal of SFT is obtain a model π^{SFT} through minimizing the cross-entropy loss:

$$\mathcal{L}_{SFT} = - \sum_{k=1}^n \log P_{\pi^0}(\hat{y}_k = y_k | \mathcal{P}, Y_{1:k-1}), \quad (7)$$

where \hat{y}_k is the k^{th} token of the generated SQL query \hat{Y} , and Y is the corresponding ground-truth label.

The SFT approach for text-to-SQL has been widely adopted in text-to-SQL research for various open-source LLMs [8, 9, 46]. Compared to in-context learning (ICL) approaches, fine-tuning paradigms are more inclined to be at a starting point in LLM-based text-to-SQL. Currently, several studies exploring a better fine-tuning method have been released. We categorize the well-designed fine-tuning methods in different groups based on their mechanisms, as shown in Tab. IV.

Enhanced Architecture: The widely-used generative pre-trained transformer (GPT) framework utilizes decoder-only transformer architecture and conventional auto-regressive decoding for text generation. Recent studies on the efficiency of LLMs have revealed a common challenge: when generating long sequences with the auto-regressive paradigm, the need to incorporate the attention mechanism results in high latency for LLMs [101, 102]. In LLM-based text-to-SQL, the speed of generating SQL queries is significantly slower compared to traditional language modeling [21, 28], which has become a challenge in constructing high-efficiency local NLIDB.

As one of the solutions, CLLMs [69] are designed to address the above challenge with an enhanced model architecture and achieve a speedup for SQL generation.

TABLE IV: Well-designed methods used in fine-tuning (FT) for LLM-based text-to-SQL. The methods in each category are ordered by release time. *The methods are utilized in multiple open-source LLMs; we select a representative model to present.

Category	Adopted by	Applied LLMs	Dataset	EX	EM	VES	Release Time	Publication Venue
Enhanced Architecture	CLLMs [69]	Deepseek*	[13]	✓			Mar-2024	ICML'24
Pre-training	CodeS [9]	StarCoder	[13, 33]	✓		✓	Feb-2024	SIGMOD'24
Data Augmentation	DAIL-SQL [71]	LLaMA*	[13, 41]	✓	✓		Aug-2023	VLDB'24
	Symbol-LLM [50]	CodeLLaMA	[13]		✓		Nov-2023	ACL'24
	CodeS [9]	StarCoder	[13, 33]	✓		✓	Feb-2024	SIGMOD'24
	StructLM [70]	CodeLLaMA	[13]		✓		Feb-2024	arXiv'24
Decomposition	DTS-SQL [71]	Mistral*	[13, 40]	✓	✓		Feb-2024	arXiv'24

Data Augmentation: During the fine-tuning process, the most straightforward factor affecting the model’s performance is the quality of the training labels [103]. The fine-tuning under the low quality or lack of the training labels is “*making bricks without straw*”, using high-quality or augmented data for fine-tuning always surpasses the meticulous design of fine-tuning methods on low-quality or raw data [29, 104]. Data-augmented fine-tuning in text-to-SQL made substantial progress, focusing on enhancing the data quality during the SFT process.

DAIL-SQL [8] are designed as an in-context learning framework, utilizing a sampling strategy for better few-shot instances. Incorporating the sampled instances in the SFT process improves the performance of open-source LLMs. Symbol-LLM [50] propose injection and infusion stage for data augmented instruction tuning. CodeS [9] augmented the training data with bi-directional generation with the help of ChatGPT. StructLM [70] are trained on multiple struct knowledge tasks for improving overall capability.

Pre-training: Pre-training is a fundamental phase of the complete fine-tuning process, aimed at acquiring text generation capabilities through auto-regressive training on extensive data [105]. Conventionally, the current powerful proprietary LLMs (e.g., ChatGPT [106], GPT-4 [78], Claude [107]) are pre-trained on hybrid corpus, which mostly benefit from the dialogue scenario that exhibits text generation capability [77]. The code-specific LLMs (e.g., CodeLLaMA [108], StarCoder [109]) are pre-trained on code data [87], and the mixture of various programming languages enables the LLMs to generate code to meet with the user’s instruction [110]. As a sub-task of code generation, the main challenge of SQL-specific pre-training technique is that the SQL/Database-related content occupies only a small portion of the entire pre-training corpus. Then, as a result, the open-source LLMs with comparatively limited comprehensive capacity (compared to ChatGPT, GPT-4) do not acquire a promising understanding of how to convert NL questions to SQL during their pre-training process.

The pre-training phase of the Codes [9] model consists of three stages of incremental pre-training. Starting from a basic code-specific LLM [109], CodeS are further pre-trained on a hybrid training corpus, including SQL-related data, NL-to-Code data, and NL-related data. The text-to-SQL understanding and performance are significantly improved.

Decomposition: Decomposing a task into multiple steps or using multiple models to solve the task is an intuitive solution for addressing a complex scenario, as we previously introduced

in Sec. IV-A, ICL paradigm. The proprietary models utilized in ICL-based methods have a massive number of parameters that are not at the same parameter level as the open-source models used in fine-tuning methods. These models inherently possess the capability to perform assigned sub-tasks well (through mechanisms such as few-shot learning) [30, 57]. Thus, to replicate the success of this paradigm in ICL methods, it is necessary to reasonably assign corresponding sub-tasks to open-source models (such as generating external knowledge, schema linking, and distilling the schema) for sub-task-specific fine-tuning and constructing the corresponding data for fine-tuning, thereby assisting in the final SQL generation.

DTS-SQL [71] proposed a two-stage decomposed text-to-SQL fine-tuning framework and designed a schema-linking pre-generation task ahead of the final SQL generation.

V. EXPECTATIONS

Despite the significant advancements made in text-to-SQL research, there are still several challenges that need to be addressed. In this section, we discuss the remaining challenges that we expect to overcome in future work.

A. Robustness in Real-world Applications

The text-to-SQL implemented by LLMs is expected to perform generalization and robustness across complex scenarios in real-world applications. Despite recent advances having made substantial progress in robustness-specific datasets [37, 41], its performance still falls short of practical application [33]. There are still challenges that are expected to be overcome in future studies. From the user aspect, there is a phenomenon that the **user is not always a clear question proposer**, which means the user questions might not have the exact database value and also can be varied from the standard datasets, the synonyms, typos, and vague expressions could be included [40]. For instance, the models are trained on clear indicative questions with concrete expressions in the fine-tuning paradigm. Since the model has not learned the mapping of realistic questions to the corresponding database, this leads to a knowledge gap when applied to real-world scenarios [33]. As reported in the corresponding evaluations of the dataset with synonym and incomplete instruction [26, 51], the SQL queries generated by ChatGPT contain around 40% incorrect execution, which is 10% lower than the original evaluation [51]. Simultaneously, the **fine-tuning with local text-to-SQL datasets may contain non-standardized samples and labels**. As an example, the name

of the table or column is not always an accurate representation of its content, which yields an inconsistency within the training data construction and may lead to a semantic gap between the database schema and the user question. To address this challenge, aligning the LLMs with intention bias and designing the training strategy towards noisy scenarios will benefit the recent advances. At the same time, **the data size in real-world applications is relatively smaller than the research-oriented benchmark**. Since extending a large amount of the data by human annotation incurs high labor costs, designing data-augmentation methods to obtain more question-SQL pairs will support the LLM in data scarcity. Also, the adaptation study of fine-tuned open-source LLM to the local small-size dataset can be potentially beneficial. Furthermore, **the extensions on multi-lingual [42, 111] and multi-modal scenarios [112]** should be studied comprehensively in future research, which will benefit more language groups and help build more general database interfaces.

B. Computational Efficiency

The computational efficiency is determined by the inference speed and the cost of computational resources, which is worth considering in both application and research work [49, 69]. With the increasing complexity of databases in up-to-date text-to-SQL benchmarks [15, 33], databases will carry more information (including more tables and columns), and the token length of the database schema will correspondingly increase, raising a series of challenges. Dealing with an ultra-complex database, taking the corresponding schema as input may encounter the challenge that **the cost of calling proprietary LLMs will significantly increase, potentially exceeding the model's maximum token length**, especially with the implementation of open-source models that have shorter context lengths. Meanwhile, another obvious challenge is that most works **use the full schema as model input, which introduces significant redundancy** [57]. Providing LLMs with a precise question-related filtered schema directly from the user end to reduce cost and redundancy is a potential solution to improve computational efficiency [30]. Designing an accurate method for schema filtering remains a future direction. Although the in-context learning paradigm achieves promising accuracy, as a computational efficiency concern, the well-designed methods with the multi-stage framework or extended context **increasing the number of API calls to enhance performance has simultaneously led to a substantial rise in costs** [7]. As reported in related approaches [49], a trade-off between performance and computational efficiency should be considered carefully, and designing a comparable (even better) in-context learning method with less API cost will be a practical implementation and is still under exploration. Compared to PLM-based methods, **the inference speed of LLM-based methods is observably slower** [21, 28]. Accelerating inference by shortening the input length and reducing the number of stages in implementation would be intuitive for the in-context learning paradigm. For local LLMs, from a starting point [69], more speedup strategies can be studied in enhancing the model's architecture in future exploration.

C. Data Privacy and Interpretability

As a part of the LLMs' study, LLM-based text-to-SQL also faces some general challenges present in LLM research [4, 113, 114]. Potential improvements from the text-to-SQL perspective are also expected to be seen in these challenges, thereby extensively benefiting the study of LLMs. As previously discussed in Sec. IV-A, the in-context learning paradigm predominates the number and performance in recent studies, with the majority of work using proprietary models for implementation [7, 8]. A straightforward challenge is proposed regarding data privacy, as **calling proprietary APIs to handle local databases with confidentiality can pose a risk of data leakage**. Using a local fine-tuning paradigm can partially address this issue. Still, the current performance of vanilla fine-tuning is not ideal [8], and advanced fine-tuning framework potentially relies on proprietary LLMs for data augmentation [9]. Based on the current status, more tailored frameworks in the local fine-tuning paradigm for text-to-SQL deserve widespread attention. Overall, the development of deep learning continually faces challenges regarding interpretability [114, 115]. As a long-standing challenge, considerable work has already been studied to address this issue [116, 117]. However, in text-to-SQL research, **the interpretability of LLM-based implementation is still not being discussed**, whether in the in-context learning or fine-tuning paradigm. The approaches with a decomposition phase explain the text-to-SQL implementation process from the perspective of step-by-step generation [7, 51]. Building on this, combining advanced study in interpretability [118, 119] to enhance text-to-SQL performance and interpreting the local model architecture from the database knowledge aspect remain future directions.

D. Extensions

As a sub-field of LLMs and natural language understanding research, many studies in these fields have been adopted for text-to-SQL tasks, advancing its development [89, 95]. However, text-to-SQL research can also be extended to the larger scope studies of these fields at meanwhile. For instance, SQL generation is a part of code generation. The well-designed approaches in code generation also obtain promising performance in text-to-SQL [48, 68], performing generalization across various programming languages. **The potential extension of some tailored text-to-SQL frameworks to NL-to-code studies can also be discussed**. For instance, frameworks integrating execution output in NL-to-code can also achieve solid performance in SQL generation [7]. An attempt to extend execution-aware approaches in text-to-SQL with other advancing modules [30, 31] for code generation is worth discussing. From another perspective, we previously discussed that text-to-SQL can enhance LLM-based question-answering (QA) by providing factual information. The database can store relational knowledge as structural information, and **the structure-based QA can potentially benefit from text-to-SQL** (e.g., knowledge-based question-answering, KBQA [120]). Construct the factual knowledge with database structure, and then incorporate the text-to-SQL system to achieve information retrieval, which can potentially assist further QA with more

accurate factual knowledge [121]. More extensions of text-to-SQL studies are expected in future work.

REFERENCES

- [1] L. Wang, B. Qin, B. Hui, B. Li, M. Yang, B. Wang, B. Li, J. Sun, F. Huang, L. Si, and Y. Li, “Proton: Probing schema linking information from pre-trained language models for text-to-sql parsing,” in *International Conference on Knowledge Discovery and Data Mining (KDD)*, 2022.
- [2] B. Qin, B. Hui, L. Wang, M. Yang, J. Li, B. Li, R. Geng, R. Cao, J. Sun, L. Si *et al.*, “A survey on text-to-sql parsing: Concepts, methods, and future directions,” *arXiv preprint arXiv:2208.13629*, 2022.
- [3] S. Xu, S. Semnani, G. Campagna, and M. Lam, “Autoqa: From databases to qa semantic parsers with only synthetic training data,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [4] Y. Zhang, Y. Li, L. Cui, D. Cai, L. Liu, T. Fu, X. Huang, E. Zhao, Y. Zhang, Y. Chen *et al.*, “Siren’s song in the ai ocean: a survey on hallucination in large language models,” *arXiv preprint arXiv:2309.01219*, 2023.
- [5] P. Manakul, A. Liusie, and M. J. Gales, “Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [6] S. Lin, J. Hilton, and O. Evans, “Truthfulqa: Measuring how models mimic human falsehoods,” in *Association for Computational Linguistics (ACL)*, 2021.
- [7] M. Pourreza and D. Rafiei, “DIN-SQL: Decomposed in-context learning of text-to-SQL with self-correction,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [8] D. Gao, H. Wang, Y. Li, X. Sun, Y. Qian, B. Ding, and J. Zhou, “Text-to-sql empowered by large language models: A benchmark evaluation,” in *International Conference on Very Large Data Bases (VLDB)*, 2024.
- [9] H. Li, J. Zhang, H. Liu, J. Fan, X. Zhang, J. Zhu, R. Wei, H. Pan, C. Li, and H. Chen, “Codes: Towards building open-source language models for text-to-sql,” *arXiv preprint arXiv:2402.16347*, 2024.
- [10] F. Li and H. V. Jagadish, “Constructing an interactive natural language interface for relational databases,” in *International Conference on Very Large Data Bases (VLDB)*, 2014.
- [11] T. Yu, C.-S. Wu, X. V. Lin, bailin wang, Y. C. Tan, X. Yang, D. Radev, richard socher, and C. Xiong, “Grappa: Grammar-augmented pre-training for table semantic parsing,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [12] T. Mahmud, K. A. Hasan, M. Ahmed, and T. H. C. Chak, “A rule based approach for nlp based query processing,” in *International Conference on Electrical Information and Communication Technologies (EICT)*, 2015.
- [13] T. Yu, R. Zhang, K. Yang, M. Yasunaga, D. Wang, Z. Li, J. Ma, I. Li, Q. Yao, S. Roman, Z. Zhang, and D. Radev, “Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2018.
- [14] V. Zhong, C. Xiong, and R. Socher, “Seq2sql: Generating structured queries from natural language using reinforcement learning,” *arXiv preprint arXiv:1709.00103*, 2017.
- [15] M. Pourreza and D. Rafiei, “Evaluating cross-domain text-to-SQL models and benchmarks,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [16] I. Sutskever, O. Vinyals, and Q. V. Le, “Sequence to sequence learning with neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- [17] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [18] B. Hui, X. Shi, R. Geng, B. Li, Y. Li, J. Sun, and X. Zhu, “Improving text-to-sql with schema dependency learning,” *arXiv preprint arXiv:2103.04399*, 2021.
- [19] D. Choi, M. C. Shin, E. Kim, and D. R. Shin, “Ryan-sql: Recursively applying sketch-based slot fillings for complex text-to-sql in cross-domain databases,” in *International Conference on Computational Linguistics (COLING)*, 2021.
- [20] P. Yin, G. Neubig, W.-t. Yih, and S. Riedel, “Tabert: Pretraining for joint understanding of textual and tabular data,” *arXiv preprint arXiv:2005.08314*, 2020.
- [21] H. Li, J. Zhang, C. Li, and H. Chen, “Resdsq: Decoupling schema linking and skeleton parsing for text-to-sql,” in *Conference on Artificial Intelligence (AAAI)*, 2023.
- [22] J. Li, B. Hui, R. Cheng, B. Qin, C. Ma, N. Huo, F. Huang, W. Du, L. Si, and Y. Li, “Graphix-t5: Mixing pre-trained transformers with graph-aware layers for text-to-sql parsing,” in *Conference on Artificial Intelligence (AAAI)*, 2023.
- [23] D. Rai, B. Wang, Y. Zhou, and Z. Yao, “Improving generalization in language model-based text-to-sql semantic parsing: Two simple semantic boundary-based techniques,” in *Association for Computational Linguistics (ACL)*, 2023.
- [24] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 2019.
- [25] Q. Lyu, K. Chakrabarti, S. Hathi, S. Kundu, J. Zhang, and Z. Chen, “Hybrid ranking network for text-to-sql,” *arXiv preprint arXiv:2008.04759*, 2020.
- [26] N. Rajkumar, R. Li, and D. Bahdanau, “Evaluating the text-to-sql capabilities of large language models,” *arXiv preprint arXiv:2204.00498*, 2022.
- [27] A. Liu, X. Hu, L. Wen, and P. S. Yu, “A comprehensive evaluation of chatgpt’s zero-shot text-to-sql capability,” *arXiv preprint arXiv:2303.13547*, 2023.
- [28] J. Yang, H. Jin, R. Tang, X. Han, Q. Feng, H. Jiang, S. Zhong, B. Yin, and X. Hu, “Harnessing the power of

- llms in practice: A survey on chatgpt and beyond,” *Transactions on Knowledge Discovery from Data (TKDD)*, 2024.
- [29] W. X. Zhao, K. Zhou, J. Li, T. Tang, X. Wang, Y. Hou, Y. Min, B. Zhang, J. Zhang, Z. Dong *et al.*, “A survey of large language models,” *arXiv preprint arXiv:2303.18223*, 2023.
- [30] X. Dong, C. Zhang, Y. Ge, Y. Mao, Y. Gao, J. Lin, D. Lou *et al.*, “C3: Zero-shot text-to-sql with chatgpt,” *arXiv preprint arXiv:2307.07306*, 2023.
- [31] Z. Hong, Z. Yuan, H. Chen, Q. Zhang, F. Huang, and X. Huang, “Knowledge-to-sql: Enhancing sql generation with data expert llm,” *arXiv preprint arXiv:2402.11517*, 2024.
- [32] Q. Zhang, J. Dong, H. Chen, W. Li, F. Huang, and X. Huang, “Structure guided large language model for sql generation,” *arXiv preprint arXiv:2402.13284*, 2024.
- [33] J. Li, B. Hui, G. QU, J. Yang, B. Li, B. Li, B. Wang, B. Qin, R. Geng, N. Huo, X. Zhou, C. Ma, G. Li, K. Chang, F. Huang, R. Cheng, and Y. Li, “Can LLM already serve as a database interface? a BIG bench for large-scale database grounded text-to-SQLs,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [34] L. Wang, A. Zhang, K. Wu, K. Sun, Z. Li, H. Wu, M. Zhang, and H. Wang, “DuSQL: A large-scale and pragmatic Chinese text-to-SQL dataset,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [35] T. Yu, R. Zhang, H. Er, S. Li, E. Xue, B. Pang, X. V. Lin, Y. C. Tan, T. Shi, Z. Li, Y. Jiang, M. Yasunaga, S. Shim, T. Chen, A. Fabbri, Z. Li, L. Chen, Y. Zhang, S. Dixit, V. Zhang, C. Xiong, R. Socher, W. Lasecki, and D. Radev, “CoSQL: A conversational text-to-SQL challenge towards cross-domain natural language interfaces to databases,” in *Empirical Methods in Natural Language Processing and International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.
- [36] C.-H. Lee, O. Polozov, and M. Richardson, “KaggleD-BQA: Realistic evaluation of text-to-SQL parsers,” in *Association for Computational Linguistics and International Joint Conference on Natural Language Processing (ACL-IJCNLP)*, 2021.
- [37] X. Pi, B. Wang, Y. Gao, J. Guo, Z. Li, and J.-G. Lou, “Towards robustness of text-to-SQL models against natural and realistic adversarial table perturbation,” in *Association for Computational Linguistics (ACL)*, 2022.
- [38] Y. Gan, X. Chen, Q. Huang, and M. Purver, “Measuring and improving compositional generalization in text-to-SQL via component alignment,” in *Findings of North American Chapter of the Association for Computational Linguistics (NAACL)*, 2022.
- [39] Y. Gan, X. Chen, and M. Purver, “Exploring underexplored limitations of cross-domain text-to-SQL generalization,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2021.
- [40] Y. Gan, X. Chen, Q. Huang, M. Purver, J. R. Woodward, J. Xie, and P. Huang, “Towards robustness of text-to-SQL models against synonym substitution,” in *Association for Computational Linguistics and International Joint Conference on Natural Language Processing (ACL-IJCNLP)*, 2021.
- [41] X. Deng, A. H. Awadallah, C. Meek, O. Polozov, H. Sun, and M. Richardson, “Structure-grounded pretraining for text-to-SQL,” in *North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 2021.
- [42] Q. Min, Y. Shi, and Y. Zhang, “A pilot study for Chinese SQL semantic parsing,” in *Empirical Methods in Natural Language Processing and International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.
- [43] T. Yu, R. Zhang, M. Yasunaga, Y. C. Tan, X. V. Lin, S. Li, H. Er, I. Li, B. Pang, T. Chen, E. Ji, S. Dixit, D. Proctor, S. Shim, J. Kraft, V. Zhang, C. Xiong, R. Socher, and D. Radev, “SPaRC: Cross-domain semantic parsing in context,” in *Association for Computational Linguistics (ACL)*, 2019.
- [44] T. Shi, C. Zhao, J. Boyd-Graber, H. Daumé III, and L. Lee, “On the potential of lexico-logical alignments for semantic parsing to SQL queries,” in *Findings of Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [45] S. Xue, C. Jiang, W. Shi, F. Cheng, K. Chen, H. Yang, Z. Zhang, J. He, H. Zhang, G. Wei, W. Zhao, F. Zhou, D. Qi, H. Yi, S. Liu, and F. Chen, “Db-gpt: Empowering database interactions with private large language models,” *arXiv preprint arXiv:2312.17449*, 2024.
- [46] B. Zhang, Y. Ye, G. Du, X. Hu, Z. Li, S. Yang, C. H. Liu, R. Zhao, Z. Li, and H. Mao, “Benchmarking the text-to-sql capability of large language models: A comprehensive evaluation,” *arXiv preprint arXiv:2403.02951*, 2024.
- [47] S. Chang and E. Fosler-Lussier, “How to prompt LLMs for text-to-SQL: A study in zero-shot, single-domain, and cross-domain settings,” in *NeurIPS 2023 Second Table Representation Learning Workshop (NeurIPS)*, 2023.
- [48] X. Chen, M. Lin, N. Schärli, and D. Zhou, “Teaching large language models to self-debug,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [49] H. Zhang, R. Cao, L. Chen, H. Xu, and K. Yu, “ACT-SQL: In-context learning for text-to-SQL with automatically-generated chain-of-thought,” in *Findings of Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [50] F. Xu, Z. Wu, Q. Sun, S. Ren, F. Yuan, S. Yuan, Q. Lin, Y. Qiao, and J. Liu, “Symbol-llm: Towards foundational symbol-centric interface for large language models,” *arXiv preprint arXiv:2311.09278*, 2024.
- [51] C.-Y. Tai, Z. Chen, T. Zhang, X. Deng, and H. Sun, “Exploring chain of thought style prompting for text-to-SQL,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [52] L. Nan, Y. Zhao, W. Zou, N. Ri, J. Tae, E. Zhang, A. Cohan, and D. Radev, “Enhancing text-to-SQL

- capabilities of large language models: A study on prompt design strategies,” in *Findings of Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [53] R. Sun, S. O. Arik, H. Nakhosht, H. Dai, R. Sinha, P. Yin, and T. Pfister, “Sql-palm: Improved large language model adaptation for text-to-sql,” *arXiv preprint arXiv:2306.00739*, 2023.
- [54] S. Chang and E. Fosler-Lussier, “Selective demonstrations for cross-domain text-to-SQL,” in *Findings of Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [55] H. Xia, F. Jiang, N. Deng, C. Wang, G. Zhao, R. Mihalcea, and Y. Zhang, “Sql-craft: Text-to-sql through interactive refinement and enhanced reasoning,” *arXiv preprint arXiv:2402.14851*, 2024.
- [56] T. Zhang, T. Yu, T. B. Hashimoto, M. Lewis, W. tau Yih, D. Fried, and S. I. Wang, “Coder reviewer reranking for code generation,” in *International Conference on Machine Learning (ICML)*, 2023.
- [57] B. Wang, C. Ren, J. Yang, X. Liang, J. Bai, L. Chai, Z. Yan, Q.-W. Zhang, D. Yin, X. Sun, and Z. Li, “Mac-sql: A multi-agent collaborative framework for text-to-sql,” *arXiv preprint arXiv:2312.11242*, 2024.
- [58] Y. Xie, X. Jin, T. Xie, M. Lin, L. Chen, C. Yu, L. Cheng, C. Zhuo, B. Hu, and Z. Li, “Decomposition for enhancing attention: Improving llm-based text-to-sql through workflow paradigm,” *arXiv preprint arXiv:2402.10671*, 2024.
- [59] Y. Fan, Z. He, T. Ren, C. Huang, Y. Jing, K. Zhang, and X. S. Wang, “Metasql: A generate-then-rank framework for natural language to sql translation,” *arXiv preprint arXiv:2402.17144*, 2024.
- [60] Z. Li, X. Wang, J. Zhao, S. Yang, G. Du, X. Hu, B. Zhang, Y. Ye, Z. Li, R. Zhao, and H. Mao, “Pet-sql: A prompt-enhanced two-stage text-to-sql framework with cross-consistency,” *arXiv preprint arXiv:2403.09732*, 2024.
- [61] T. Ren, Y. Fan, Z. He, R. Huang, J. Dai, C. Huang, Y. Jing, K. Zhang, Y. Yang, and X. S. Wang, “Purple: Making a large language model a better sql writer,” *arXiv preprint arXiv:2403.20014*, 2024.
- [62] C. Guo, Z. Tian, J. Tang, P. Wang, Z. Wen, K. Yang, and T. Wang, “Prompting gpt-3.5 for text-to-sql with de-semanticization and skeleton retrieval,” in *Pacific Rim International Conference on Artificial Intelligence (PRICAI)*, 2024.
- [63] J. Jiang, K. Zhou, Z. Dong, K. Ye, X. Zhao, and J.-R. Wen, “StructGPT: A general framework for large language model to reason over structured data,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [64] C. Guo, Z. Tian, J. Tang, S. Li, Z. Wen, K. Wang, and T. Wang, “Retrieval-augmented gpt-3.5-based text-to-sql framework with sample-aware prompting and dynamic revision chain,” in *International Conference on Neural Information Processing (ICONIP)*, 2024.
- [65] D. Wang, L. Dou, X. Zhang, Q. Zhu, and W. Che, “Improving demonstration diversity by human-free fusing for text-to-sql,” *arXiv preprint arXiv:2402.10663*, 2024.
- [66] Y. Gu, Y. Shu, H. Yu, X. Liu, Y. Dong, J. Tang, J. Srinivasa, H. Latapie, and Y. Su, “Middleware for llms: Tools are instrumental for language agents in complex environments,” *arXiv preprint arXiv:2402.14672*, 2024.
- [67] F. Shi, D. Fried, M. Ghazvininejad, L. Zettlemoyer, and S. I. Wang, “Natural language to code translation with execution,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [68] A. Ni, S. Iyer, D. Radev, V. Stoyanov, W.-t. Yih, S. I. Wang, and X. V. Lin, “Lever: Learning to verify language-to-code generation with execution,” in *International Conference on Machine Learning (ICML)*, 2023.
- [69] S. Kou, L. Hu, Z. He, Z. Deng, and H. Zhang, “Cllms: Consistency large language models,” *arXiv preprint arXiv:2403.00835*, 2024.
- [70] A. Zhuang, G. Zhang, T. Zheng, X. Du, J. Wang, W. Ren, S. W. Huang, J. Fu, X. Yue, and W. Chen, “Structlm: Towards building generalist models for structured knowledge grounding,” *arXiv preprint arXiv:2402.16671*, 2024.
- [71] M. Pourreza and D. Rafiei, “Dts-sql: Decomposed text-to-sql with small large language models,” *arXiv preprint arXiv:2402.01117*, 2024.
- [72] D. Xu, W. Chen, W. Peng, C. Zhang, T. Xu, X. Zhao, X. Wu, Y. Zheng, and E. Chen, “Large language models for generative information extraction: A survey,” *arXiv preprint arXiv:2312.17617*, 2023.
- [73] G. Katsogiannis-Meimarakis and G. Koutrika, “A survey on deep learning approaches for text-to-sql,” *The VLDB Journal*, 2023.
- [74] J. Guo, Z. Zhan, Y. Gao, Y. Xiao, J.-G. Lou, T. Liu, and D. Zhang, “Towards complex text-to-sql in cross-domain database with intermediate representation,” *arXiv preprint arXiv:1905.08205*, 2019.
- [75] X. Xu, C. Liu, and D. Song, “Sqlnet: Generating structured queries from natural language without reinforcement learning,” *arXiv preprint arXiv:1711.04436*, 2017.
- [76] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “Roberta: A robustly optimized bert pretraining approach,” *arXiv preprint arXiv:1907.11692*, 2019.
- [77] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [78] J. Achiam, S. Adler, S. Agarwal, L. Ahmad, I. Akkaya, F. L. Aleman, D. Almeida, J. Altschmidt, S. Altman, S. Anadkat *et al.*, “Gpt-4 technical report,” *arXiv preprint arXiv:2303.08774*, 2023.
- [79] J. Wang, E. Shi, S. Yu, Z. Wu, C. Ma, H. Dai, Q. Yang, Y. Kang, J. Wu, H. Hu *et al.*, “Prompt engineering

- for healthcare: Methodologies and applications,” *arXiv preprint arXiv:2304.14670*, 2023.
- [80] J. Wei, M. Bosma, V. Y. Zhao, K. Guu, A. W. Yu, B. Lester, N. Du, A. M. Dai, and Q. V. Le, “Finetuned language models are zero-shot learners,” *arXiv preprint arXiv:2109.01652*, 2021.
 - [81] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar *et al.*, “Llama: Open and efficient foundation language models,” *arXiv preprint arXiv:2302.13971*, 2023.
 - [82] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei, N. Bashlykov, S. Batra, P. Bhargava, S. Bhosale *et al.*, “Llama 2: Open foundation and fine-tuned chat models,” *arXiv preprint arXiv:2307.09288*, 2023.
 - [83] J. Bai, S. Bai, Y. Chu, Z. Cui, K. Dang, X. Deng, Y. Fan, W. Ge, Y. Han, F. Huang *et al.*, “Qwen technical report,” *arXiv preprint arXiv:2309.16609*, 2023.
 - [84] L. Reynolds and K. McDonell, “Prompt programming for large language models: Beyond the few-shot paradigm,” in *ACM Conference on Human Factors in Computing Systems (CHI)*, 2021.
 - [85] X. Ye and G. Durrett, “The unreliability of explanations in few-shot prompting for textual reasoning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
 - [86] C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu, “Exploring the limits of transfer learning with a unified text-to-text transformer,” *The Journal of Machine Learning Research (JMLR)*, 2020.
 - [87] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. d. O. Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman *et al.*, “Evaluating large language models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
 - [88] J. Zamfirescu-Pereira, R. Y. Wong, B. Hartmann, and Q. Yang, “Why johnny can’t prompt: how non-ai experts try (and fail) to design llm prompts,” in *ACM Conference on Human Factors in Computing Systems (CHI)*, 2023.
 - [89] J. Wei, X. Wang, D. Schuurmans, M. Bosma, F. Xia, E. Chi, Q. V. Le, D. Zhou *et al.*, “Chain-of-thought prompting elicits reasoning in large language models,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
 - [90] D. Zhou, N. Schärli, L. Hou, J. Wei, N. Scales, X. Wang, D. Schuurmans, C. Cui, O. Bousquet, Q. Le *et al.*, “Least-to-most prompting enables complex reasoning in large language models,” *arXiv preprint arXiv:2205.10625*, 2022.
 - [91] W. Lei, W. Wang, Z. Ma, T. Gan, W. Lu, M.-Y. Kan, and T.-S. Chua, “Re-examining the role of schema linking in text-to-SQL,” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
 - [92] Q. Liu, D. Yang, J. Zhang, J. Guo, B. Zhou, and J.-G. Lou, “Awakening latent grounding from pretrained language models for semantic parsing,” in *Findings of Association for Computational Linguistics (ACL)*, 2021.
 - [93] Z. Tan, X. Liu, Q. Shu, X. Li, C. Wan, D. Liu, Q. Wan, and G. Liao, “Enhancing text-to-SQL capabilities of large language models through tailored promptings,” in *International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)*, 2024.
 - [94] J. Huang and K. C.-C. Chang, “Towards reasoning in large language models: A survey,” in *Findings of Association for Computational Linguistics (ACL)*, 2023.
 - [95] X. Wang, J. Wei, D. Schuurmans, Q. V. Le, E. H. Chi, S. Narang, A. Chowdhery, and D. Zhou, “Self-consistency improves chain of thought reasoning in language models,” in *International Conference on Learning Representations (ICLR)*, 2023.
 - [96] W. Chen, X. Ma, X. Wang, and W. W. Cohen, “Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks,” *Transactions on Machine Learning Research (TMLR)*, 2023.
 - [97] M. Müller and R. Sennrich, “Understanding the properties of minimum Bayes risk decoding in neural machine translation,” in *Association for Computational Linguistics and International Joint Conference on Natural Language Processing (ACL-IJCNLP)*, 2021.
 - [98] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn, “Direct preference optimization: Your language model is secretly a reward model,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
 - [99] G. Team, T. Mesnard, C. Hardin, R. Dadashi, S. Bhupatiraju, S. Pathak, L. Sifre, M. Rivière, M. S. Kale, J. Love *et al.*, “Gemma: Open models based on gemini research and technology,” *arXiv preprint arXiv:2403.08295*, 2024.
 - [100] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray *et al.*, “Training language models to follow instructions with human feedback,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
 - [101] Y. Leviathan, M. Kalman, and Y. Matias, “Fast inference from transformers via speculative decoding,” in *International Conference on Machine Learning (ICML)*, 2023.
 - [102] C. Chen, S. Borgeaud, G. Irving, J.-B. Lespiau, L. Sifre, and J. Jumper, “Accelerating large language model decoding with speculative sampling,” *arXiv preprint arXiv:2302.01318*, 2023.
 - [103] H. Song, M. Kim, D. Park, Y. Shin, and J.-G. Lee, “Learning from noisy labels with deep neural networks: A survey,” *IEEE Transactions on Neural Networks and Learning Systems (TNNLS)*, 2023.
 - [104] N. Deng, Y. Chen, and Y. Zhang, “Recent advances in text-to-SQL: A survey of what we have and what we expect,” in *International Conference on Computational Linguistics (COLING)*, 2022.
 - [105] D. Erhan, A. Courville, Y. Bengio, and P. Vincent, “Why does unsupervised pre-training help deep learning?” in *Artificial Intelligence and Statistics (AISTATS)*, 2010.
 - [106] Y. Liu, T. Han, S. Ma, J. Zhang, Y. Yang, J. Tian, H. He, A. Li, M. He, Z. Liu *et al.*, “Summary of chatgpt-related research and perspective towards the future of large

- language models,” *Meta-Radiology*, 2023.
- [107] Anthropic, “Introducing Claude,” 2023.
- [108] B. Roziere, J. Gehring, F. Gloeckle, S. Sootla, I. Gat, X. E. Tan, Y. Adi, J. Liu, T. Remez, J. Rapin *et al.*, “Code llama: Open foundation models for code,” *arXiv preprint arXiv:2308.12950*, 2023.
- [109] R. Li, L. B. allal, Y. Zi, N. Muennighoff, D. Kocetkov, C. Mou, M. Marone, C. Akiki, J. LI, J. Chim, Q. Liu, E. Zheltonozhskii, T. Y. Zhuo, T. Wang, O. Dehaene, J. Lamy-Poirier, J. Monteiro, N. Gontier, M.-H. Yee, L. K. Umapathi, J. Zhu, B. Lipkin, M. Oblokulov, Z. Wang, R. Murthy, J. T. Stillerman, S. S. Patel, D. Abulkhanov, M. Zocca, M. Dey, Z. Zhang, U. Bhattacharyya, W. Yu, S. Luccioni, P. Villegas, F. Zhdanov, T. Lee, N. Timor, J. Ding, C. S. Schlesinger, H. Schoelkopf, J. Ebert, T. Dao, M. Mishra, A. Gu, C. J. Anderson, B. Dolan-Gavitt, D. Contractor, S. Reddy, D. Fried, D. Bahdanau, Y. Jernite, C. M. Ferrandis, S. Hughes, T. Wolf, A. Guha, L. V. Werra, and H. de Vries, “Starcoder: may the source be with you!” *Transactions on Machine Learning Research (TMLR)*, 2023.
- [110] Y. Wang, W. Zhong, L. Li, F. Mi, X. Zeng, W. Huang, L. Shang, X. Jiang, and Q. Liu, “Aligning large language models with human: A survey,” *arXiv preprint arXiv:2307.12966*, 2023.
- [111] A. Tuan Nguyen, M. H. Dao, and D. Q. Nguyen, “A pilot study of text-to-SQL semantic parsing for Vietnamese,” in *Findings of Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [112] Y. Song, R. C.-W. Wong, and X. Zhao, “Speech-to-sql: toward speech-driven sql query generation from natural language question,” *The VLDB Journal*, 2024.
- [113] B. Yan, K. Li, M. Xu, Y. Dong, Y. Zhang, Z. Ren, and X. Cheng, “On protecting the data privacy of large language models (llms): A survey,” *arXiv preprint arXiv:2403.05156*, 2024.
- [114] C. Singh, J. P. Inala, M. Galley, R. Caruana, and J. Gao, “Rethinking interpretability in the era of large language models,” *arXiv preprint arXiv:2402.01761*, 2024.
- [115] D. Dai, L. Dong, Y. Hao, Z. Sui, B. Chang, and F. Wei, “Knowledge neurons in pretrained transformers,” in *Association for Computational Linguistics (ACL)*, 2022.
- [116] N. Zhang, Y. Yao, B. Tian, P. Wang, S. Deng, M. Wang, Z. Xi, S. Mao, J. Zhang, Y. Ni *et al.*, “A comprehensive study of knowledge editing for large language models,” *arXiv preprint arXiv:2401.01286*, 2024.
- [117] K. Meng, A. S. Sharma, A. J. Andonian, Y. Belinkov, and D. Bau, “Mass-editing memory in a transformer,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [118] K. Meng, D. Bau, A. Andonian, and Y. Belinkov, “Locating and editing factual associations in gpt,” *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [119] C. Zheng, L. Li, Q. Dong, Y. Fan, Z. Wu, J. Xu, and B. Chang, “Can we edit factual knowledge by in-context learning?” in *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [120] H. Luo, Z. Tang, S. Peng, Y. Guo, W. Zhang, C. Ma, G. Dong, M. Song, W. Lin *et al.*, “Chatkbqa: A generate-then-retrieve framework for knowledge base question answering with fine-tuned large language models,” *arXiv preprint arXiv:2310.08975*, 2023.
- [121] G. Xiong, J. Bao, and W. Zhao, “Interactive-kbqa: Multi-turn interactions for knowledge base question answering with large language models,” *arXiv preprint arXiv:2402.15131*, 2024.
- [122] R. Anil, A. M. Dai, O. Firat, M. Johnson, D. Lepikhin, A. Passos, S. Shakeri, E. Taropa, P. Bailey, Z. Chen *et al.*, “Palm 2 technical report,” *arXiv preprint arXiv:2305.10403*, 2023.
- [123] Z. Hong and J. Liu, “Towards better question generation in qa-based event extraction,” *arXiv preprint arXiv:2405.10517*, 2024.
- [124] Y. Liu, H. He, T. Han, X. Zhang, M. Liu, J. Tian, Y. Zhang, J. Wang, X. Gao, T. Zhong *et al.*, “Understanding llms: A comprehensive overview from training to inference,” *arXiv preprint arXiv:2401.02038*, 2024.
- [125] A. Zeng, X. Liu, Z. Du, Z. Wang, H. Lai, M. Ding, Z. Yang, Y. Xu, W. Zheng, X. Xia, W. L. Tam, Z. Ma, Y. Xue, J. Zhai, W. Chen, Z. Liu, P. Zhang, Y. Dong, and J. Tang, “GLM-130b: An open bilingual pre-trained model,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [126] Q. Zhang, J. Dong, Q. Tan, and X. Huang, “Integrating entity attributes for error-aware knowledge graph embedding,” *IEEE Transactions on Knowledge and Data Engineering (TKDE)*, 2024.
- [127] Q. Zhang, J. Dong, H. Chen, X. Huang, D. Zha, and Z. Yu, “Knowgpt: Black-box knowledge injection for large language models,” *arXiv preprint arXiv:2312.06185*, 2023.
- [128] F. Huang, Z. Yang, J. Jiang, Y. Bei, Y. Zhang, and H. Chen, “Large language model interaction simulator for cold-start item recommendation,” *arXiv preprint arXiv:2402.09176*, 2024.
- [129] Y. Bei, H. Xu, S. Zhou, H. Chi, M. Zhang, Z. Li, and J. Bu, “Cpdg: A contrastive pre-training method for dynamic graph neural networks,” *arXiv preprint arXiv:2307.02813*, 2023.
- [130] Y. Bei, H. Chen, S. Chen, X. Huang, S. Zhou, and F. Huang, “Non-recursive cluster-scale graph interacted model for click-through rate prediction,” in *International Conference on Information and Knowledge Management (CIKM)*, 2023.
- [131] Z. Yuan, D. Liu, W. Pan, and Z. Ming, “Sql-rank++: A novel listwise approach for collaborative ranking with implicit feedback,” in *International Joint Conference on Neural Networks (IJCNN)*, 2022.
- [132] H. Chen, Y. Bei, Q. Shen, Y. Xu, S. Zhou, W. Huang, F. Huang, S. Wang, and X. Huang, “Macro graph neural networks for online billion-scale recommender systems,” in *International World Wide Web Conference (WWW)*, 2024.